FlashNS: Operator- and Adjoint-Aware Contracted Derivative Execution for High-Order PINNs

Dezhen Lu
University of Tübingen
dezhen.lu@student.uni-tuebingen.de

Updated research draft | October 13, 2026

Abstract

High-order physics-informed training requires coordinate derivatives and their complete parameter gradients. FlashNS executes consumer-specific derivative states with parameter transposes and bounded chunk workspaces. In the original two-dimensional FP64 Navier-Stokes experiment on NVIDIA H20, all 18 representation-study solves and all 54 matched public-comparison solves pass unchanged numerical and independent scalar-AD acceptance. The eight-channel GL8 implementation gives 1.659× paired setup-plus-optimization speedup over compiled compact Q10; the public collapsed Taylor, explicit differentiation and exact two-dimensional STDE adaptations give respective baseline/GL8 ratios of 1.730, 2.312 and 3.368. We extend the study to a three-dimensional transient curl-potential model with a 36-channel partial derivative state, 4,999 objective points and all 50,692 parameters. Full H100 objective-gradient checks pass the frozen numerical budget. A packed-gather candidate nevertheless remains slower than C36 Taylor-state parameter AD: paired baseline/candidate ratios are 0.973 in eager mode and 0.970 with CUDA Graph replay. A separately compiled packed activation and analytic activation transpose give a 1.265× paired objective-gradient speedup on one CPU thread (95% bootstrap interval 1.203-1.328); this variant has no CUDA qualification or GPU performance result. Both H100 training engines complete their frozen two-seed budgets, but neither passes the prescribed convergence and physical acceptance criteria. The results establish the stated execution tradeoffs and numerical qualifications; they do not establish three-dimensional time-to-accuracy acceleration or a fluid regularity theorem.

1 Introduction

Training a physics-informed neural network (PINN) requires derivatives of a network output with respect to coordinates and derivatives of the resulting loss with respect to every network parameter. These are different consumers. For gradient-enhanced incompressible Navier-Stokes training, the residual uses second coordinate derivatives and its gradient uses third derivatives. A fast residual evaluation does not establish a fast training step, and a fast training step does not establish a shorter solve to a prescribed error.

Taylor-mode and explicit derivative methods already address the cost of nested coordinate automatic differentiation (AD). Residual-specific partial jets, collapsed Taylor propagation and stochastic Taylor contractions are especially relevant to this problem [3, 4, 5]. We build on that setting. Our question is whether a concrete operator contraction, its complete parameter transpose and a practical bounded working-state schedule can improve an unchanged third-order training objective beyond a compiled full-jet implementation and public derivative methods.

For this objective, independent third partials are not needed: the diffusion consumer asks for the gradient of the Laplacian. We retain the scalar value, coordinate gradient, symmetric Hessian and gradient of the Laplacian. In two dimensions this is an eight-channel state, GL8, instead of the ten channels of a complete third-order state, Q10. Closure through an activation includes products between lower derivatives. Its

parameter transpose additionally depends on the fourth activation derivative. Omitting that dependence would change the parameter gradient even if some forward residuals agreed.

FlashNS makes this consumer contract explicit and implements both sides of it. It keeps matrix multiplications in library GEMMs, uses explicit CUDA rules for pointwise derivative propagation and residual seeds, and replays the complete loss/gradient computation with current parameter storage. Boundary points use a value-only path. A chunked mode reuses one interior workspace and accumulates the complete globally weighted parameter gradient before any optimizer update.

We evaluate the implementation on the original FP64 Kovasznay gradient-enhanced objective, including all six numerical acceptance criteria and a separate scalar-AD final check. The representation study uses six initialization seeds; the public comparison reuses those planned seeds and reports its actual coverage. Paired comparisons use identical data and initial parameter bytes. We report full-gradient replay, setup, optimization and process wall time separately. Public competing implementations are pinned and adapted to the same objective; their adaptations and application limits are stated explicitly. Development qualification never supplies a formal solve result.

The resulting contribution is a verified execution case study: a closed contracted derivative state with its complete parameter VJP, a reusable chunked implementation, and a source-pinned comparison of actual solves. It does not establish the first residual-aware derivative method, a globally minimal state, or superiority across PDEs, precisions or GPU families.

This revision preserves the October 6 H20 campaign and adds a separate three-dimensional transient C36 study in Section 6. The two campaigns have different models, source data, parameter counts, seeds, acceptance gates and hardware; their timings are never pooled. The H20 comparisons remain accepted solve measurements. The new H100 and CPU ratios measure complete objective/gradient evaluations, while the new three-dimensional solver acceptance remains unsuccessful.

2 Contracted derivative execution

2.1 A fixed scientific objective

We study steady incompressible flow with constant viscosity ν and physical Cartesian coordinates $x = (x_1, x_2)$. A neural network predicts (u, v, p). Its residual is

$$\begin{aligned} r_u &= u u_{x_1} + v u_{x_2} + p_{x_1} - \nu \Delta u, \\ r_v &= u v_{x_1} + v v_{x_2} + p_{x_2} - \nu \Delta v, \\ r_d &= u_{x_1} + v_{x_2}. \end{aligned} \tag{1}$$

Following the gradient-enhanced objective [2], we retain both the residual and its coordinate gradient:

$$L = \tfrac{1}{2}\mathrm{mean}_\Omega \left(\|r\|_2^2 + 0.2\|\nabla r\|_F^2 \right) + 5\left(\mathrm{mean}_{\partial\Omega} \|(u, v) - (u_*, v_*)\|_2^2 + \mathrm{mean}_{x_1 = 1}(p - p_*)^2 \right). \tag{2}$$

This objective needs third coordinate derivatives. Removing the gradient penalty would change the scientific problem. All implementations use the existing point weights and validation definitions; in particular, the boundary component RMS is not the training vector-square mean.

2.2 A closed representation

For a scalar intermediate z, define

$$S(z) = (z, g, H, t), \qquad g = \nabla z, \quad H = \nabla^2 z, \quad t = \nabla \Delta z. \tag{3}$$

Only independent symmetric Hessian entries are stored. In two dimensions the raw, unnormalized state is

$$(z, z_{x_1}, z_{x_2}, z_{x_1x_1}, z_{x_1x_2}, z_{x_2x_2}, z_{x_1x_1x_1} + z_{x_1x_2x_2}, z_{x_1x_1x_2} + z_{x_2x_2x_2}). \tag{4}$$

We call this state GL8. It has eight channels per hidden scalar, compared with ten for the complete third-order two-dimensional Taylor state Q10. The raw-to-Taylor projection includes the appropriate factorials; GL8 channels must not be read as Taylor coefficients.

For a pointwise activation $y = \phi(z)$, let $a_k = \phi^{(k)}(z)$. Direct application of the chain rule gives

$$\begin{aligned} g_y &= a_1 g, \\ H_y &= a_1 H + a_2 g g^\top, \\ t_y &= a_1 t + a_2(g \operatorname{tr} H + 2Hg) + a_3 g \|g\|_2^2. \end{aligned} \tag{5}$$

An affine layer applies the same weight matrix to each state channel, with the bias added only to the value channel. Consequently, affine layers and smooth pointwise activations form a closed execution rule for S. The state supplies every quantity in r and ∇r: convection gradients need the Hessian, while diffusion gradients need t.

Closure is sufficient for this consumer, not a claim of global minimality or a reconstruction of all third derivatives. For example, $z = x_1^3 - 3x_1x_2^2$ has $S(z) = 0$ at the origin but $z_{x_1x_1x_1} = 6$. A request for an independent third partial cannot be answered from GL8. The runtime rejects unsupported coordinate identities or contractions instead of returning such an observable.

2.3 Complete parameter adjoints

Reducing the forward state alone is insufficient for training. FlashNS transposes the same local closed rules, including their dependence on the primal value. To make that dependence explicit, define $Q = g \operatorname{tr} H + 2Hg$ and $c = g^\top g$. The differential of the third contracted component is

$$\begin{aligned} \delta t_y = {} & a_1 \delta t + a_2(\delta g \operatorname{tr} H + g \operatorname{tr} \delta H + 2\delta H g + 2H\delta g) \\ & + a_3\big(c\,\delta g + 2g(g^\top \delta g)\big) + (a_2 t + a_3 Q + a_4 g c)\delta z. \end{aligned} \tag{6}$$

Thus the local reverse rule requires a_4, although the forward coordinate state stops at order three. The value, gradient and Hessian differentials are transposed together with (6). For the packed symmetric Hessian, off-diagonal entries use the transpose of symmetric unpacking: the two matrix contributions are added, not treated as two independent stored channels.

If X is a layer input state and $\bar{Z}$ its preactivation adjoint, the complete parameter updates are

$$\bar{W} = \sum_{n,q} \bar{Z}_{nq} X_{nq}^\top, \qquad \bar{b} = \sum_n \bar{Z}_{n,0}, \qquad \bar{X}_{nq} = W^\top \bar{Z}_{nq}. \tag{7}$$

The implementation uses flattened point/channel matrix multiplications for these reductions. A fixed polynomial residual program seeds all nine residual and residual-gradient terms. It is a specific, verified consumer rule rather than a general symbolic residual compiler.

In exact arithmetic, projection of a complete coordinate jet commutes with the closed affine/activation rules. Composing with a loss that depends only on that projection therefore preserves its parameter derivative by the ordinary chain rule. Our numerical checks test both the projected forward quantities and the entire flattened parameter gradient; forward agreement alone would not establish this property.

2.4 Stable activation and bounded working state

The production kernels reuse the existing stable tanh convention. In particular, the first derivative is evaluated with a stable exponential tail instead of relying exclusively on subtraction from a rounded saturated value. Higher derivatives and their reverse rules use the same primal convention. Affine work remains in library GEMMs; pointwise activation and residual transposes use explicit CUDA kernels. CUDA Graph replay reads the current parameter storage rather than a snapshot, which is checked by consecutive in-place updates.

Dirichlet and outlet consumers that need only values use a separate value-only path. The chunked implementation processes interior points in blocks of 512 and accumulates the complete parameter gradient using the original global weights. It never normalizes each chunk separately. Its working-state bound is proportional to chunk size, while fixed data, parameters and optimizer storage remain outside that bound. Floating reduction order can alter line-search trajectories; every backend must still satisfy the same independent final criteria.

The measured GL8 production scope is two-dimensional, constant-viscosity, pointwise affine/tanh networks with first parameter VJPs. The current contract does not support a learned contraction, nonlinear coordinate changes, sample-coupled layers, nonlocal losses, HVPs or double backward.

3 Relation to existing derivative methods

Physics-informed networks put PDE constraints into a training objective [1]; gradient-enhanced training also penalizes residual derivatives [2]. Our contribution concerns execution of that unchanged derivative-intensive objective.

Cao et al. [3] already exploit residual-dependent partial jets and specialized layer rules, checking residuals and parameter gradients. FlashNS therefore does not claim priority for propagating only PDE-relevant derivative information. The comparison here examines a specific contracted state together with an explicit complete parameter transpose and bounded interior working storage.

Dangel et al. [4] collapse derivative sums within Taylor-mode computation for linear differential operators. This is closely related to retaining contractions instead of a full tensor. We evaluate the public collapsed transform, including parameter backward and the additional residual-gradient consumer. Merely comparing with our own older compiled full-jet implementation would not constitute this external baseline.

STDE and its extension [6, 5] use high-order Taylor directions and contractions to avoid materializing large derivative tensors. The JMLR code link currently points to the public STDE repository. Our two-dimensional exact-coordinate adaptation evaluates that code; it is neither a separate implementation of every STDE++ mixed-partial construction nor an experiment on high-dimensional stochastic amortization. Those distinctions bound the comparison.

The baseline adapters and precise source commits are part of the artifact. Their measured results pertain to the matched scientific objective and execution modes, not to each method's full application range or the speedups originally reported on other workloads.

4 Experimental protocol

4.1 Problem, data and acceptance

The analytic Kovasznay solution is

$$u_* = 1 - e^{\lambda x_1}\cos(2\pi x_2), \quad v_* = \frac{\lambda}{2\pi}e^{\lambda x_1}\sin(2\pi x_2), \quad p_* = \frac{1-e^{2\lambda x_1}}{2}, \qquad \lambda = \frac{1}{2\nu} - \sqrt{\frac{1}{4\nu^2} + 4\pi^2}. \tag{8}$$

We set $\nu = 0.07$ on $[-0.5, 1] \times [-0.5, 1.5]$. The network has widths 2-64-64-3, two tanh hidden layers, a linear head and 4,547 parameters. Training uses 2,048 interior points and 128 points on each of four edges. The 128 pressure-outlet points share the last edge's coordinates; there are 512 boundary coordinates, not 640. We retain the original globally normalized interior, velocity and outlet weights.

The fixed stop-check set contains 8,192 interior points and 257 points on each boundary edge. A separate scalar nested-AD interior set verifies every candidate final solution. Both validations must meet all six limits in Table 1, with finite values. PDE RMS combines the three residuals by the original sum-of-squares convention. Boundary velocity RMS averages velocity components, whereas the training loss uses their vector square. We preserve these definitions rather than reconstructing them from metric names. Five data/weight SHA256 values are identical across every matched run; initialization affects parameters only.

Table 1. Unchanged numerical acceptance limits. A run succeeds only if both the fixed validation and independent scalar-AD validation satisfy every row.

Metric	Upper limit
Combined PDE residual RMS	0.02
Boundary velocity component RMS	0.01
Outlet pressure RMS	0.01
Relative L_2 error of u	0.01
Relative L_2 error of v	0.02
Relative L_2 error of p	0.02

4.2 Optimizer and finite attempts

Every formal arm uses the same optimizer and stopping program. Adam permits 4,000 updates with learning rate 10^{-3}, halved every 1,000 updates, $\beta = (0.9, 0.999)$ and $\varepsilon = 10^{-8}$. Stop checks occur before training and every 100 Adam updates. Strong-Wolfe L-BFGS then permits at most 500 blocks, each with 20 iterations, nominal 25 evaluations, history size 50, learning rate 1, gradient tolerance 10^{-12} and change tolerance 10^{-15}. Stop checks follow each block. Line searches may exceed a nominal evaluation limit; reported closure counts are actual counts.

The two-dimensional formal seeds are 862101-862106. Its development seeds 361901 and 361902 and older regression seeds are excluded from the two-dimensional formal comparisons. The new three-dimensional campaign uses its separately frozen original seeds 361901 and 361902. No seed, threshold or optimizer budget is changed after an observed result. Every attempt is retained. Nonconvergence consumes its measured cost and counts against the attempted arm; an unsupported execution variant does not turn into a successful baseline through substitution. Paired time ratios are computed only when both members pass and are accompanied by all outcomes.

The Q10/GL8/chunk campaign balances the six permutations of three arms. The subsequent public comparison has six paper-group arms; three sampled permutations and their reverses balance their order over the same six seeds. NVIDIA baselines use separately admitted windows, so their timings are not interleaved with paper-group controls. This shared H20 runs one admitted GPU window at a time. All recorded children finish naturally. Whole-window ownership includes copying and checking evidence before explicit handback.

4.3 Baselines and execution modes

Q10 compact propagates the complete two-dimensional third-order Taylor state at interior points and values at boundaries. It uses the existing explicit parameter transpose, full-graph Torch compilation and CUDA Graph replay. It is a strong internal representation control, not the public collapsed Taylor package.

Cao ED calls the pinned author's `ed_u_jet` without changing that function. Its value, first derivatives and diagonal second derivatives assemble the Navier-Stokes residual. Coordinate AD of that residual adds the original gradient-enhancement term; parameter AD supplies the complete gradient. This is a matched-objective adaptation of the public implementation. We also execute the unmodified original author NS training recipe separately; its architecture, objective and timing are not pooled with matched solves.

Collapsed Taylor calls the public `jet(collapsed=True)` transform with exact coordinate directions. The collapsed second-order output supplies the Laplacian, and residual coordinate AD supplies the additional third-order consumer. Parameter AD is included. We test both the default complete loss/gradient program and its traced, compiled variant.

STDE exact calls the public `hess_diag` sparse-STDE operator with dimension two, `rand_batch_size=0` and no randomized Jacobian. The two coordinate directions are exhausted, so this experiment uses no stochastic derivative estimator. JAX forward residual differentiation and parameter backward preserve the same loss. JAX runs on the GPU in FP64; CPU fallback is rejected. DLPack exchange, copies and synchronization with the common Torch optimizer are included in replay and solve wall time. This experiment does not evaluate every STDE++ mixed-partial construction or its intended high-dimensional stochastic setting.

Torch AD evaluates nested coordinate AD and complete parameter backward on an ordinary dense MLP. **PhysicsNeMo** [8] uses the public `PhysicsInformer` with symbolic Navier-Stokes equations and its autodiff gradient method, followed by residual-gradient AD and parameter backward. It is the full residual path, not just a standalone gradient component.

cuDNN [9] uses an algebraically equivalent 1×1 convolution MLP with coordinate and parameter AD. A cuDNN operator must appear in the actual profiler trace; a fallback is retained as a failed qualification rather than called a cuDNN result. **cuEquivariance polynomial** [10] reuses the inherited FP64 Q10 segmented-polynomial forward and explicit parameter VJP with the native `uniform_1d` backend. That implementation retains its unmodified all-points Q10 policy, including boundaries. Scalar-irreps `Linear` fused-tensor-product variants are separate diagnostics and cannot replace the polynomial formal arm. Actual native operators are required in the profiler.

For Torch-based public adapters and exact STDE, high-order derivatives run only at the 2,048 interior points and values only at 512 boundary points, matching the Q10/GL8 controls' point policy. Public adapters use `torch.tanh`, while production FlashNS preserves its stable tail convention. Mathematically they represent the same activation; numerical full-gradient equivalence is checked rather than assumed.

4.4 Qualification, caches and timing scopes

Qualification evaluates the complete loss and all 4,547 parameter gradients against the full-jet reference after each of three consecutive in-place updates. The unchanged numerical tolerance is absolute 2×10^{-12} and relative 2×10^{-11}. It also checks actual native execution where required. Passing default/compiled variants are selected by the smallest development median complete-gradient wall time; the selection map and report hashes are frozen before formal seeds. No failed variant is used.

Qualification warms task-owned Torch, native-extension, Triton and JAX caches. Each formal attempt is a fresh process with new parameters, optimizer state and graphs. Persistent JAX compilation caching is explicitly enabled, with compile threshold zero and entry threshold -1; cache settings and hit/miss logs are retained. The resulting solve experiment is a warm-dependency execution comparison. Initial qualification/build cost is reported separately and is not silently amortized over the six seeds.

Complete-gradient qualification uses five groups of 30 replays after a two-second warm-up and reports the median group wall time. The earlier representation/memory study instead uses five independently ordered processes per shape/arm, each with seven groups of 100 replays; these are separate experiments. Replay timing includes full loss, every parameter gradient and chunk/interoperation costs, but excludes optimizer, stop validation and setup.

Optimization wall time includes training, optimizer work, line searches, synchronization, stop checks and independent final acceptance. Setup plus optimization additionally includes per-run engine preparation. It excludes interpreter startup, common data creation and teardown; full child-process wall time includes those costs. We report paired geometric means of time ratios, not ratios of unpaired medians. Small-sample ranges are descriptive, not a general-PDE confidence guarantee.

The GPU is NVIDIA H20 (SM90, 97,871 MiB), driver 580.76.05, CUDA toolkit 13.0.88, Torch 2.12.1+cu130 and Triton 3.7.1. All matched experiments use FP64 and TF32 is disabled. JAX is 0.11.2; the native cuEquivariance ops package is 0.12.0. New dependencies live in a task-owned environment inheriting the unchanged base numerical stack. This original two-dimensional campaign reports H20 measurements only. The separate H100 and CPU measurements added in this revision are described in Section 6; no new Blackwell, SM120 or multi-GPU result is claimed.

The representation, memory and profiler study used the earlier physical H20 `edf64e5f`; the current public comparison, with six planned seeds, uses physical H20 `d952768b`. All current public measurements use the latter device; the planned NVIDIA comparisons use that same device. Earlier public campaigns remain archived separately and are not pooled with the current public timing comparison. Pairwise ratios are computed within their own device-specific campaign.

5 Two-dimensional H20 results

5.1 Complete-gradient qualification

Table 2. Development GPU qualification on seed 361902. Median complete loss/all-parameter-gradient wall time from five groups of 30 replays; optimizer, validation and setup excluded. Each passing variant satisfies three successive updated-parameter checks. JAX exchange is included. A missing value indicates failed qualification, not zero time.

Method	Selected execution	Gradient ms	Max gradient error
FlashNS GL8	Native + graph	1.583	7.55×10^{-15}
Q10 compact	Inductor + graph	3.267	3.55×10^{-15}
Collapsed Taylor	Compiled + graph	3.751	5.33×10^{-15}
Cao ED	Compiled + graph	5.725	5.33×10^{-15}
Torch coordinate AD	Compiled + graph	7.651	5.33×10^{-15}
STDE exact (2D)	JAX JIT	9.459	5.33×10^{-15}
PhysicsNeMo	Compiled + graph	6.506	5.33×10^{-15}
cuDNN convolution AD	Compiled + graph	74.004	5.33×10^{-15}
cuEquivariance polynomial	Native polynomial + graph	3.340	5.33×10^{-15}

9 of 9 paper-group development variants pass the complete-gradient checks; their largest observed absolute parameter-gradient discrepancy is 7.55×10^{-15}. Selected variants are Q10 compact (q10_compact-default), FlashNS GL8 (gl8_native-default), Torch coordinate AD (torch_ad-compiled), Cao ED (ed_author-compiled), Collapsed Taylor (torchjet_collapsed-compiled), STDE exact (2D) (stde_exact-default). Selection uses development replay latency only, never observed formal solve speed or convergence. Qualification costs and every alternative remain in the artifact.

The separate NVIDIA qualification attempts 9 variants, of which 5 pass. Failed native-path or complete-gradient checks remain in the raw reports and consume their measured process cost. Only selected passing formal arms are trained; unsupported planned cells remain not-run.

The cuEquivariance polynomial adapter requires an initialization-only portability repair: its unused inherited FlashNS native-build initialization is omitted, while polynomial evaluation and all derivative and parameter-gradient functions remain inherited without modification. An initial recovery attempt fails at import before numerical evaluation and is retained with its process cost. The corrected recovery passes all three updated-parameter checks and the native-path profile before formal training. Scalar tensor-product variants fail higher-order autograd qualification; these failures characterize the tested route and do not establish a general limitation of cuEquivariance.

5.2 Six-seed representation and chunk comparison

Table 3. Representation campaign: all 18 attempts pass fixed and independent acceptance. Time columns are medians in seconds. Last column is the geometric mean of six paired Q10/candidate ratios for setup plus optimization, not the ratio of table medians.

Method	Accepted	Optimization	Setup + opt.	Process	Paired speedup
Q10 compact	6/6	33.660	38.519	41.798	1.000×
FlashNS GL8	6/6	22.756	23.110	25.598	1.659×
GL8 chunk-512	6/6	28.004	28.354	31.091	1.423×

Native GL8 is faster than Q10 in all six paired solves. The optimization-only geometric mean is 1.479×, while setup plus optimization is 1.659× (pair range 1.460-1.809). Full-process speedup is 1.635×. Chunk-512 has setup-plus-optimization speedup 1.423×. Its six-seed median is longer than native GL8,

consistent with the extra replay and accumulation schedule. Development preflight cost is 24.962 seconds and is reported separately.

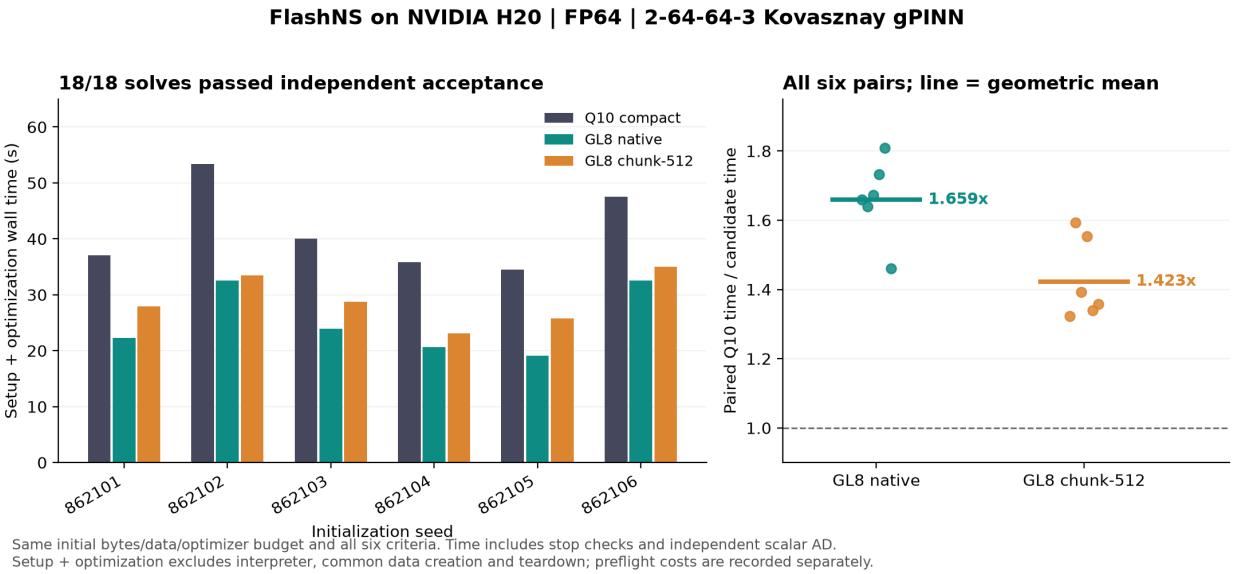


Figure 1. Every observed paired solve in the representation campaign. All 18 pass. Connected points identify backends across different initialization seeds; they do not interpolate a continuous experimental variable.

5.3 Source-pinned public solves

Table 4. Public comparison. Outcomes are accepted/nonconverged/runtime-or-correctness-failure/not-run out of six planned seeds. Time medians in seconds retain all attempted solves. Final column is paired baseline/GL8 setup-plus-optimization geometric mean and number of jointly accepted pairs; values above one favor GL8. NVIDIA arms run in separate windows.

Method	Outcomes	Opt. s	Setup + opt. s	Process s	Ratio / n
FlashNS GL8	6/0/0/0	23.131	23.496	25.818	1.000 / 6
Q10 compact	6/0/0/0	34.090	39.248	42.877	1.658 / 6
Collapsed Taylor	6/0/0/0	35.955	40.415	43.775	1.730 / 6
Cao ED	6/0/0/0	50.055	54.810	58.232	2.312 / 6
Torch coordinate AD	6/0/0/0	64.800	70.682	74.447	2.994 / 6
STDE exact (2D)	6/0/0/0	75.376	78.650	81.400	3.368 / 6
PhysicsNeMo	6/0/0/0	54.479	72.506	77.286	3.108 / 6
cuDNN convolution AD	6/0/0/0	523.492	531.241	534.903	22.865 / 6
cuEquivariance polynomial	6/0/0/0	32.155	35.729	38.697	1.555 / 6

The public matrix plans 54 cells and records 54 actual formal child attempts, of which 54 pass every acceptance criterion. The priority paper group records 36 of 36 planned attempts. All observed formal initial parameter hashes match the corresponding seeds in the original representation campaign; the six planned initializations are reused across experiments, not additional independent seeds for each backend.

Paired baseline/GL8 setup-plus-optimization ratios are Collapsed Taylor: 1.730× (6 accepted pairs, range 1.482-1.906); Cao ED: 2.312× (6 accepted pairs, range 2.024-2.607); STDE exact (2D): 3.368× (6 accepted

pairs, range 2.833-3.862); Torch coordinate AD: 2.994× (6 accepted pairs, range 2.712-3.259). They describe the matched adaptations and common optimizer pipeline. They are not the original papers' speedup claims.

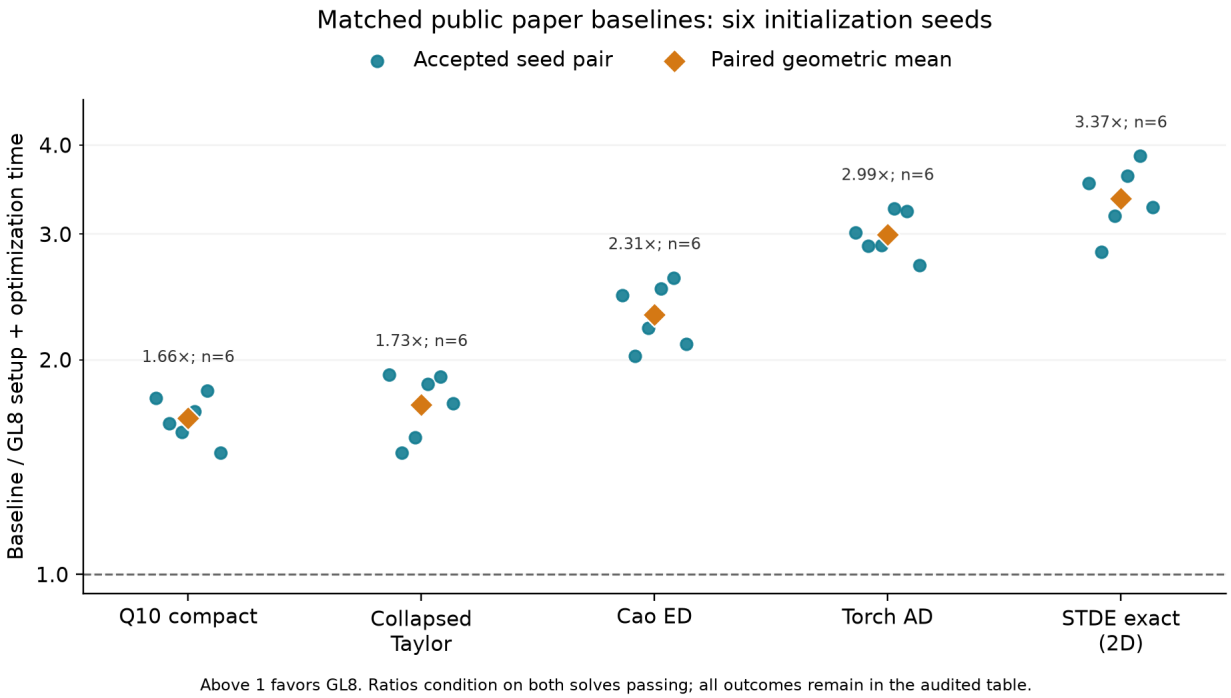


Figure 2. Paper-group solve ratios for every jointly accepted initialization pair. Above one favors GL8; the dashed line is equal time. Horizontal offsets separate seed observations, not additional experimental conditions. Table 4 retains all attempted outcomes and the accepted-pair denominators.

Table 5. Largest final error over all accepted seeds and both validators, rounded to four decimals. Acceptance uses the unrounded values and every criterion; n is the number of accepted seeds. Unsupported or unattempted methods have no final-error statistic.

Method	n	**PDE**	**Boundary**	**Outlet**	u	v	p
FlashNS GL8	6	0.0138	0.0093	0.0025	0.0095	0.0200	0.0099
Q10 compact	6	0.0147	0.0096	0.0024	0.0096	0.0199	0.0109
Collapsed Taylor	6	0.0132	0.0096	0.0025	0.0093	0.0200	0.0091
Cao ED	6	0.0136	0.0095	0.0026	0.0091	0.0200	0.0103
Torch coordinate AD	6	0.0146	0.0099	0.0028	0.0097	0.0199	0.0102
STDE exact (2D)	6	0.0141	0.0097	0.0030	0.0094	0.0200	0.0096
PhysicsNeMo	6	0.0146	0.0095	0.0021	0.0096	0.0199	0.0097
cuDNN convolution AD	6	0.0157	0.0100	0.0023	0.0093	0.0200	0.0112
cuEquivariance polynomial	6	0.0150	0.0098	0.0024	0.0094	0.0199	0.0097

The nvidia qualification retains 4 failed execution variants. Their complete tracebacks and process costs are in the raw artifact. Formal cells with no passing variant remain not-run; no failed backend is renamed or replaced by the scalar diagnostic.

Qualification process-wall costs, including unsuccessful variants and the original author recipe where applicable, are nvidia 137.927 s, papers 263.687 s. These costs are separate from warm-dependency formal solve timings.

5.4 Memory and replay tradeoff

Table 6. Earlier independent-process replay/memory study: five processes per shape/arm, seven groups of 100 replays per process. Latency is the median of process medians. Allocator measurements are taken before allocating the full-jet correctness oracle; they measure Torch allocation/reservation, not total GPU process memory.

Interior	Method	Gradient ms	Setup peak MiB	Ready reserved MiB
2,048	Q10 compact	3.272	87.16	102.00
2,048	FlashNS GL8	1.578	69.61	76.00
2,048	GL8 chunk-512	2.271	42.92	56.00
16,384	Q10 compact	23.539	456.72	642.00
16,384	FlashNS GL8	11.511	312.79	340.00
16,384	GL8 chunk-512	17.908	45.88	76.00

At 16,384 interior points, chunk-512 reduces native GL8 setup/capture allocated peak from 312.79 to 45.88 MiB, a reduction of 85.33%, and ready reserved memory from 340 to 76 MiB. Full-gradient replay increases from 11.511 to 17.908 ms, a 55.58% cost. At 2,048 points, native GL8 remains the speed-oriented option. No maximum trainable scale or device-traffic reduction is inferred from these counters. A separate 2,048-point native profile attributes 82.30% of kernel-time sum to library GEMMs; that single-profile share is not a solve-wall attribution.

5.5 Original author recipe sanity check

The unmodified Cao NS script naturally completes on the public-comparison GPU. It uses its own six-hidden-layer 64-wide network, FP32 lid-driven-cavity objective and three 1,000-Adam repeats. The archived summary gives AD and ED mean times of 11.132 and 4.613 seconds respectively (ratio 2.413). Its complete-gradient consistency report is retained. This separate reproduction confirms the public script runs; it is not a Kovasznay gPINN solve or a FlashNS comparison.

6 Three-dimensional transient C36 study

6.1 Scientific problem and derivative contract

This separate manufactured transient incompressible-flow experiment uses the frozen physical-coordinate source cache and calibration, viscosity $\nu = 0.01$, cutoff width $\eta = 1/32$, and the affine/tanh network widths 4-128-128-128-128-4. Its outputs represent a three-component curl potential and pressure, with scales 0.01 and 0.1 respectively. The prescribed moving-box geometry supplies input features and a compact gate; the model and particle right-hand side use the learned potential, without inserting the exact interior velocity. The curl construction enforces divergence freedom in real arithmetic. Numerical divergence is still measured rather than assumed to establish all physical acceptance.

For a scalar potential intermediate, the retained raw derivative state is

$$S_{36}(z) = (z, g, H, T, Q, m, n, K), \qquad Q_{ij} = \partial_{ij}\Delta z, \tag{9}$$

where g, H and T are the spatial gradient, symmetric Hessian and symmetric third derivative; $m = \partial_t z$, $n = \nabla \partial_t z$ and $K = \nabla^2 \partial_t z$. These blocks have $1 + 3 + 6 + 10 + 6 + 1 + 3 + 6 = 36$ independent channels. They supply the curl velocity, transient momentum residual and its gradient, including contracted fourth spatial derivatives of the potential. C36 is a partial-jet contract for this consumer, not GL8, a full fourth-order jet, or a minimality result.

The canonical objective contains 4,096 interior, 512 initial, 384 boundary and seven pressure-gauge points. All 4,999 points and all 50,692 parameters are included. Frozen physical normalization and the original global family denominators are retained, including during chunk accumulation. The source and calibration hashes and complete scientific protocol are identified in the artifact [11]. FP64 and disabled TF32 are fixed. This campaign's two original stored initializations are 361901 and 361902; they are not extra H20 formal seeds.

6.2 Execution changes and numerical qualification

The H100 packed-gather revision transposes symmetric unpacking through four cached integer gathers instead of a chain of 54 in-place slice assignments. It retains the same C36 algebra and local state-autograd nonlinear pullbacks. Its full-objective GPU proof records 48 comparisons and 96 saved arrays, covering both original and changed parameters. The maximum discrepancy is 2.07227×10^{-5} of the unchanged $10^{-11} + 10^{-9}|g_{\mathrm{ref}}|$ comparison budget. Its complete native reference derives from the earlier all-point H100 qualification, which freshly evaluated native coordinate and parameter AD at every canonical point; the gather qualification does not rerun that expensive native reference. Partition agreement alone is not an independent coordinate oracle.

The subsequent opt-in packed activation computes the six Hessian and ten third-derivative entries directly, avoiding their full 9- and 27-entry expansion and repacking. A manual analytic transpose reverses the activation's packed contractions and the actual stable coefficient program. Gate and curl nonlinear pullbacks still use local state autograd; this is not a fully manual nonlinear adjoint for the entire model. For the activation primal $y = \tanh z$ and $e = \exp(-2|z|)$, the frozen coefficient convention is

$$\begin{aligned} a_1 &= 4e/(1+e)^2, & a_2 &= -2ya_1, \\ a_3 &= 4y^2 a_1 - 2a_1^2, & a_4 &= 16ya_1^2 - 8y^3 a_1. \end{aligned} \tag{10}$$

The transpose differentiates these evaluated operations, including the rounded tanh derivative and the exponential tail, rather than substituting an idealized fifth derivative at saturated values. Regression checks include values 0, $\pm 10^{-16}$, $\pm 10^{-12}$, ± 20 and ± 40, with cotangents scaled to 10^8.

Full canonical CPU qualification at original and changed parameters records 42 comparisons and 84 preserved arrays for the compiled analytic variant. A fresh NumPy-only reader passes every comparison; its maximum error is 9.17612×10^{-6} of the original budget. This rechecks the preserved H100 native reference ancestry rather than freshly running native coordinate AD on all CPU points. Exact initial tensor bytes come from the original checkpoints: regeneration under the newer Torch version was rejected after 1,804 of 50,692 entries differed by one ULP. That failed attempt is retained. No tolerance or scientific data was changed. Ten CPU regression checks pass. These facts do not qualify this new variant's CUDA execution, CUDA Graph replay, training or physical acceptance.

6.3 Paired complete-objective performance

Each reported ratio is the geometric mean of matched baseline/candidate mean evaluation times. Two frozen initializations each have five alternating AB/BA pairs, for ten pairs per comparison. Candidate and C36 parameter-AD windows last at least 60 seconds and contain at least three complete evaluations. Setup, compilation, warm-up and graph capture are excluded from these ratios and reported separately. CUDA timings synchronize each complete evaluation. The native-coordinate-AD leg has its original special rule of one complete native evaluation per window. The 95% intervals use 10,000 paired bootstrap resamples with seed 795912; they describe these timing pairs, not uncertainty across PDEs, datasets, hardware or independent training seeds.

Table 7. Complete C36 objective and all-parameter-gradient timing against the internal C36 parameter-AD control. Times are means over paired windows in seconds; the ratio is a paired geometric mean, not the ratio of these time means. Ten pairs per row. H100 rows use the packed-gather source; the CPU row uses a later compiled packed activation with analytic activation transpose and an equally compiled forward in its parameter-AD baseline. These different versions and devices are not pooled.

Device / mode	Candidate s	AD s	Ratio	95% interval
H100 eager	1.739582	1.692531	0.972935	0.970474-0.975924
H100 graph	0.484235	0.469551	0.969678	0.969317-0.970075
CPU compiled	1.865900	2.352072	1.265086	1.203149-1.327900

The H100 source `ae51ca643280` remains 2.782% slower in eager mode and 3.127% slower when both arms replay CUDA Graphs. The parameter-AD control shares C36 forward algebra and is a strong execution baseline; it is not an independent coordinate derivative oracle or the public collapsed Taylor library used in the H20 campaign. GPU timing and saved gradient readers complete naturally with exit zero.

The newer CPU variant uses Torch 2.14.1, NumPy 2.5.4, CPU Inductor and one Torch thread. Both arms compile the same packed forward; only the candidate supplies the analytic activation transpose. A fresh NumPy-only reader checks 24 saved gradient arrays, 12 gradient comparisons and all 20 paired timing windows. All checks pass at the original budget. The 1.265× result is CPU objective-gradient throughput; there is no GPU speed measurement or accepted solve for this variant.

A separate older source, `73baa3b281fb`, measures native coordinate AD with its original recursive derivative cache. Mean candidate/native times are 2.777272/3,484.283522 seconds, with paired ratio 1,254.839 and interval 1,211.337-1,291.263 (ten pairs). That candidate predates packed gather, and the baseline predates the later cache-lifetime repair. The large ratio is specific to that slow recursive native path; it is not a speedup against the optimized parameter-AD control, the repaired native implementation or a complete accepted solver. The timing worker, reader and timing driver exit zero. The original whole controller exits one due to its earlier failed training stage, which remains a failed controller outcome rather than being relabeled successful.

6.4 Frozen-budget training outcomes

Both H100 engines start from their original fresh initializations and execute 4,000 Adam updates plus 20 strong-Wolfe L-BFGS calls per seed. Adam has initial learning rate 10^{-3}, halved every 1,000 updates, betas (0.9,0.999) and epsilon 10^{-8}. L-BFGS uses learning rate one, history 20, at most 20 iterations and nominal 25 evaluations per call, zero gradient/change stopping tolerances and a 520-closure cap per seed. The complete-gradient checks run at frozen Adam checkpoints and after every L-BFGS call. Each engine passes all 48 scheduled saved-gradient comparisons. Successful execution and gradient checks do not establish convergence: the original joint gate requires loss at most 10^{-4} and gradient infinity norm at most 10^{-3}.

Table 8. All four C36 training attempts finish the original finite budget. Actual L-BFGS closure counts, final loss and full-gradient infinity norm are shown. Every attempt fails the joint convergence gate; no additional optimizer budget is silently appended.

Engine	Seed	Closures	Loss	$\|g\|_\infty$	Gate
Explicit	361901	444	0.00054021	0.00075071	Fail
Explicit	361902	438	0.00051678	0.00128521	Fail
Parameter AD	361901	449	0.00053870	0.00059044	Fail
Parameter AD	361902	443	0.00051532	0.00206306	Fail

All four losses remain above 10^{-4}. The explicit seed 361902 and parameter-AD seed 361902 also exceed the gradient limit. The evaluated training sources are `7281faed9ec4` (explicit) and `de1d223` (parameter AD); they do not train the later CPU analytic activation variant. No accepted three-dimensional time-to-accuracy or whole-solver acceleration is available.

6.5 Independent physical observations and failures

The unchanged holdout seed 794701 supplies 4,096 stratified and 4,096 uniform-volume points, seven time slices with 256 points each, 512 initial points, 384 boundary points and seven pressure-gauge points. Native coordinate AD observes the final learned model on H100; NumPy reconstructs the physical metrics. Learned-velocity trajectories use SciPy DOP853 for 27 particles at 129 uniform output times plus the critical times. Both original refinements complete: $(\mathrm{rtol}, \mathrm{atol}, \max \mathrm{step})$ are $(10^{-9}, 10^{-11}, 1/64)$ and $(10^{-11}, 10^{-13}, 1/128)$. The exact velocity is used for error comparison, never as the trajectory right-hand side. Both stratified and uniform field gates and the trajectory gate must pass; time-slice and regional reports remain diagnostics.

Table 9. C36 final field failures. S/U denote stratified/uniform holdouts. Velocity errors are relative L2; momentum and corrected pressure use the original physical normalization. Row limits are velocity 0.01, momentum 0.01 and pressure 0.001. All displayed values exceed their corresponding gate; complete region and time-slice values remain in the result JSON.

Engine	Seed	Vel. S	Vel. U	Mom. S	Press. S	Press. U
Explicit	361901	0.020528	0.043743	0.014220	0.025935	0.003022
Explicit	361902	0.018808	0.034605	0.013856	0.022480	0.003407
Parameter AD	361901	0.020331	0.043126	0.014276	0.024746	0.002837
Parameter AD	361902	0.018497	0.034154	0.013832	0.019825	0.003094

Table 10. C36 trajectories from the stricter original DOP853 refinement. Endpoint error must be at most 0.001; refinement change must be at most 10^{-7}. All refinement tests pass but all endpoint tests fail, so no row passes physical acceptance.

Engine	Seed	Endpoint	Refinement change	Gate
Explicit	361901	0.023166	7.2297×10^{-11}	Fail
Explicit	361902	0.029803	7.2587×10^{-11}	Fail
Parameter AD	361901	0.022762	7.2325×10^{-11}	Fail
Parameter AD	361902	0.029980	7.1729×10^{-11}	Fail

For all four checkpoints, initial/boundary velocity and gauge errors are exactly zero. Stratified normalized divergence remains at most 2.01×10^{-14}, and divergence-gradient RMS at most 3.05×10^{-12}; their limits are 10^{-3} and 0.02. Momentum-gradient RMS also passes its 0.02 limit. These satisfied gates coexist with failed velocity, stratified momentum, corrected-pressure and trajectory-endpoint criteria. Accurate divergence and tightly converged ODE integration do not establish an accurate learned velocity field.

The observer source `1f1983a2b015` repairs a recursive native derivative closure cycle that retained CUDA graphs until garbage collection. The unchanged checkpoints and physical protocol then complete; the earlier OOM and its raw logs remain archived. The repair preserves derivative order and parameter dependence. It is checked against original native states and all 50,692 gradient entries at two points; sixteen repeated native observations keep allocated CUDA memory constant at 67,514,880 bytes. This limited memory qualification supplies no repaired-native paired performance result.

Each checkpoint's remote and local NumPy readers pass 117 saved-array checks. They verify hashes, loss/gradient anchors, field residual metrics, trajectory error summaries and refinement statistics. They do not rerun the optimizer trajectory or independently reintegrate the ODE. The reader checks analytic forcing at its original tolerance and retains sealed source bytes for cancellation-sensitive residual subtraction; a prior cross-version rounding failure is preserved and no tolerance is widened.

6.6 Evidence and current limits

The public machine-readable results preserve every timing pair, all four training outcomes, and all final physical observations, with full source/protocol/raw-array hashes [11]. GPU timing and physical archives have fully verified cluster and local copies; the complete training histories have verified cluster archives, while smaller training anchors are also available locally. The complete new CPU packet is verified locally only. Publication of compact result JSON files is not publication of every raw checkpoint or profiler trace.

The unmodified source cache and calibration remain frozen. Neither optimization budgets, physical limits, accepted seeds nor global loss denominators are weakened in response to these failures. The new analytic activation is an opt-in CPU-qualified research path. Its CUDA qualification, fair paired GPU comparison, trained accuracy and accepted solve timing remain open work; no such result is implied by compilation or CPU speed.

7 Discussion and limits

The derivative consumer offers an execution opportunity even after replacing nested coordinate AD with Taylor-mode computation. GL8 removes independent third-order information that the fixed diffusion-gradient consumer never reads, while preserving lower derivatives required by nonlinear convection. Its explicit transpose is part of the optimization program, not an afterthought to a forward-only operator. The measured effect nevertheless combines representation, pointwise kernels, graph scheduling and parameter backward; the native-versus-compiled-Q10 solve comparison does not isolate a pure representation effect.

Replay and solve ratios need not match. Optimizer work, line searches and fixed validation remain outside the accelerated derivative program. Small floating-point differences can change an L-BFGS trajectory and its closure count. Setup also differs between compiled public adapters, JAX and explicit native kernels, even with warmed dependency caches. We therefore report every timing scope and actual iteration/evaluation count rather than extrapolating a replay factor to the whole training run.

Chunking addresses a different engineering constraint. It bounds the intermediate execution workspace and permits a substantial allocator reduction at the measured larger shape, while paying for input copies, extra replay launches and gradient accumulation. The fixed data and optimizer remain resident. We measured neither an OOM boundary nor the maximum trainable point count, and do not infer those from the two tested shapes. Torch allocator counters cannot rank total process memory against JAX's allocator.

The external comparisons must be read within their adaptations. ED and collapsed Taylor supply public second-order operators extended to a third-order gradient-enhanced consumer; this is not their original training benchmark. Exact two-coordinate STDE preserves a deterministic target but does not test stochastic amortization in high dimensions. The unmodified cuEquivariance polynomial arm has an all-points Q10 boundary policy, unlike the compact interior/value-only policy of the paper adapters. Its timing therefore reflects both a library implementation and that inherited policy. These differences are recorded rather than erased by a common name.

The representation campaign's six initializations provide a paired view of one fixed-data manufactured PDE. The public comparison reports its actual completed-seed coverage separately. These initializations do not measure variation across PDEs, sampled datasets, network families or hardware. Repeated replay groups within a qualification process are timing observations, not additional independent training seeds. The prior representation study's independent processes are kept separate from the public qualification process. No speed claim in this paper is extended to stochastic STDE++ constructions, variable viscosity, nonlinear coordinate maps, sample-coupled layers, second parameter derivatives or nonlocal losses.

The new three-dimensional transient case adds full-objective derivative qualification, paired execution measurements and complete failed physical observations. It has not established an accepted three-dimensional solve. A stronger general-purpose evaluation would add a second nonlinear PDE, accepted transient solutions, variable-coefficient operators and a broader activation/network family. It would also compare cold start and native framework optimizers, measure actual device traffic and examine generalized consumer selection. These remain open experiments rather than implied consequences of the current implementation.

8 Conclusion

FlashNS gives a concrete consumer contract for high-order PINN training: propagate a closed value/gradient/Hessian/gradient-Laplacian state and transpose every local dependency into the complete parameter gradient. The H20 experiments evaluate that contract under an unchanged FP64 Navier-Stokes objective and independent numerical acceptance. The compact native path and its chunked alternative offer distinct speed and workspace tradeoffs. The original H20 evidence supports that tested solve setting. The added C36 study qualifies full derivatives but finds no H100 advantage over the strong parameter-AD control, and its completed training budgets fail physical acceptance. The later CPU analytic activation achieves a paired 1.265× objective-gradient speedup with an equally compiled forward baseline. GPU transfer of that improvement and accepted three-dimensional solve acceleration remain unestablished; broader scientific claims require further experiments.

References

[1] M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686-707, 2019. https://doi.org/10.1016/j.jcp.2018.10.045.

[2] Jeremy Yu, Lu Lu, Xuhui Meng, and George Em Karniadakis. Gradient-enhanced physics-informed neural networks for forward and inverse PDE problems. arXiv:2111.02801, 2021. https://arxiv.org/abs/2111.02801.

[3] Wenbo Cao, Zhe Lu, and Weiwei Zhang. Verified residual-specific explicit derivative kernels for physics-informed learning and discretized PDE adjoints. arXiv:2606.29702, 2026. https://arxiv.org/abs/2606.29702.

[4] Felix Dangel, Tim Siebert, Marius Zeinhofer, and Andrea Walther. Collapsing Taylor Mode Automatic Differentiation. *Advances in Neural Information Processing Systems*, 38, 2025. https://papers.nips.cc/paper_files/paper/2025/hash/d835e25415cfe4d6d48a0be3b7b763bb-Abstract-Conference.html.

[5] Zekun Shi, Zheyuan Hu, Min Lin, and Kenji Kawaguchi. STDE++: Polynomial-Time Amortization for Linear Differential Operators. *Journal of Machine Learning Research*, 27(114):1-50, 2026. https://www.jmlr.org/papers/v27/25-1474.html.

[6] Zekun Shi, Zheyuan Hu, Min Lin, and Kenji Kawaguchi. Stochastic Taylor Derivative Estimator: Efficient amortization for arbitrary differential operators. *Advances in Neural Information Processing Systems*, 37, 2024. https://proceedings.neurips.cc/paper_files/paper/2024/hash/dd2eb5250696753ea37141bbd89bb569-Abstract-Conference.html.

[7] sail-sg. STDE official public implementation. Source commit `fae88663b1d2f1b0666d4d250c38a3c34b852ac0`. https://github.com/sail-sg/stde.

[8] NVIDIA. PhysicsNeMo public implementation. Source commit `b45a5c810c741e6b41f8515be24c51121f8fc21f`. https://github.com/NVIDIA/physicsnemo.

[9] NVIDIA. CUDA Deep Neural Network library (cuDNN). https://developer.nvidia.com/cuda/cuda-x-libraries/cudnn.

[10] NVIDIA. cuEquivariance public implementation. Source commit `70674ada7a00372e88ebe99969c700a6dcddd08f`. https://github.com/NVIDIA/cuEquivariance.

[11] Dezhen Lu. FlashNS C36 experiment and reproducibility artifact. Evidence snapshot `1634120bc27869ea62f5d2439d5e565ae23a58b2`, October 13, 2026; per-experiment runtime hashes are retained in the result files. Public C36 source, frozen protocols and result JSON files.

A Explicit local transpose

For completeness, let $\bar{y}$, $\bar{g}_y$, B_y and $\bar{t}_y$ denote output adjoints. Here B_y is symmetric and represents the Frobenius pairing with the full symmetric Hessian. If the stored off-diagonal Hessian channel has adjoint β_{12}, this convention sets $(B_y)_{12} = (B_y)_{21} = \beta_{12}/2$. Define $h = \operatorname{tr} H$, $c = g^\top g$ and $Q = gh + 2Hg$. Transposing the local differential gives

$$\begin{aligned}
\bar{t} &= a_1 \bar{t}_y, \\
\bar{H} &= a_1 B_y + a_2((\bar{t}_y^\top g) I + \bar{t}_y g^\top + g \bar{t}_y^\top), \\
\bar{g} &= a_1 \bar{g}_y + 2a_2 B_y g + a_2(h\bar{t}_y + 2H\bar{t}_y) + a_3(c\bar{t}_y + 2g(g^\top \bar{t}_y)), \\
\bar{z} &= a_1 \bar{y} + a_2 \bar{g}_y^\top g + \langle B_y, a_2 H + a_3 g g^\top \rangle_F + \bar{t}_y^\top (a_2 t + a_3 Q + a_4 g c).
\end{aligned} \tag{11}$$

The packed input Hessian adjoint is obtained by transposing symmetric unpacking: diagonal entries are unchanged and each off-diagonal channel receives the sum of its two matrix entries. This convention prevents a factor-of-two error. All terms contribute to the shared parameter VJP through the affine transpose; no activation derivative order is discarded because it exceeds the forward coordinate order.

The argument assumes fixed physical coordinates, coordinate-independent affine parameters, a fixed contraction and sufficiently smooth pointwise activations. Let P map a complete third-order raw jet to (z, g, H, t). The chain rule establishes $PF_{\text{full}} = F_{\text{GL8}}P$ for each supported layer. Induction gives the same identity for a supported network. Since the residual loss factors through P, its real-arithmetic parameter derivative agrees with the full-jet derivative. This is a closure and transpose argument for a stated consumer; it is not a proof that no other representation could use fewer channels.

B Chunk accumulation and storage contract

Partition interior and boundary points into finite chunks. With the original weights, the loss is the sum of chunk contributions and its parameter gradient is the sum of their parameter gradients. The executor zeros accumulators once at the beginning of a closure, copies a chunk into the reusable graph inputs, replays, and accumulates its loss and gradient. Tail padding has zero weight. Every chunk uses the same parameter version; the optimizer can update parameters only after the complete closure.

For maximum chunk size M, intermediate derivative storage is proportional to M times the retained layer-channel widths. Fixed point data still grow with total point count, and parameter/optimizer storage is not bounded by M. The eight-channel contraction removes two hidden coordinate channels; it does not by itself prove a 20% reduction in total live memory. At 2,048 interior points, compact logical point/channel rows decrease from $10(2048) + 512 = 20,992$ to $8(2048) + 512 = 16,896$. Actual allocator measurements include further temporaries and graph pools.

Captured outputs are borrowed until the next replay. Graph parameter storage can be updated in place; replacing its storage or using an incompatible stream/layout is rejected. Comparing a previously returned borrowed scalar after another closure would compare an overwritten value. The validation harness checks each live loss/gradient before replay and checks updated parameter storage separately.

C Original H20 source and evidence provenance

The original two-dimensional H20 execution source is commit `286d7e394c8bd39f9aa409715afe9f1bac915cec`, based on public main `7c261ea5635cdcec1bf98b72d1105dcf0b486087` and a preserved local implementation snapshot. The baseline adapters are separate frozen runtime files. Table 11 identifies upstream sources; full tree/blob hashes, runtime file SHA256 values and all invocation arguments accompany the artifact. Pinning distinguishes the public collapsed transform from the project's own compiled full-jet backend.

Table 11. Pinned public baseline sources. Full commit IDs are retained in the machine-readable source manifests.

Implementation	Public repository	Commit prefix
Cao ED	Cao-WenBo/ExplicitDifferentiation	`a524575baa62`
Collapsed Taylor	f-dangel/torch-jet	`2e49005ee4ae`
STDE	sail-sg/stde	`fae88663b1d2`
PhysicsNeMo	NVIDIA/physicsnemo	`b45a5c810c74`
cuEquivariance	NVIDIA/cuEquivariance	`70674ada7a00`

Original JSON reports, progress histories, complete profiler event names, stdout/stderr, final parameter tensors, protocols and natural process-exit receipts are preserved. Every archive is copied off the GPU instance and checked against its recorded whole-archive SHA256 before extraction; every manifested member is then checked individually. Analysis rejects altered data, optimizer configurations, unmatched initial parameters, repeated seed/arm attempts and claimed convergence that violates either validator. Archive-level verification includes metadata as well as numerical reports.

One inherited optimizer-protocol description still lists the earlier `formal_initialization_seeds` 75101-75103. That field does not select parameters in the solver. The frozen campaign seed list, each actual child's `-seed` argument and the reported initial tensor-byte hashes identify the six formal initializations 862101-862106. We preserve the original metadata bytes and state this discrepancy rather than silently rewriting the raw reports.

Preparation-only failures are also retained. A local boundary-cardinality assertion initially counted shared outlet coordinates twice; the frozen problem always had 512 boundary coordinates, and no numerical data changed. An earlier archive check correctly rejected an SCP transfer before transfer completion. Neither event caused a formal training retry or a changed scientific threshold. GPU comparison uses only the final frozen compact adapter packet, whose 460 runtime source hashes are checked before each admitted window.

The initial implementation regression ran 547 CPU checks with two existing upstream Euler dependency skips and 129 full-width GPU checks. Full-width Compute Sanitizer memcheck completed 112 checks with zero errors. Small network racecheck, initcheck and synccheck each completed 60 checks with zero errors; an earlier full-width racecheck was interrupted for instrumentation cost and is not counted as completed. These tests establish stated regression coverage, not correctness for every unsupported consumer.

D Every planned public solve

Table 12. All 54 planned public cells, without failed-attempt filtering. Setup plus optimization is in seconds; closure counts are actual. Six paired initialization seeds are reused across methods.

Seed	Method	Outcome	Setup + opt.	Closures
862101	FlashNS GL8	accepted	21.854	6833
862102	FlashNS GL8	accepted	32.786	9032
862103	FlashNS GL8	accepted	25.138	7279
862104	FlashNS GL8	accepted	21.068	6635
862105	FlashNS GL8	accepted	18.966	6261
862106	FlashNS GL8	accepted	30.989	8685
862101	Q10 compact	accepted	38.649	6733
862102	Q10 compact	accepted	53.417	9089
862103	Q10 compact	accepted	39.846	7076
862104	Q10 compact	accepted	35.679	6481
862105	Q10 compact	accepted	34.340	6309
862106	Q10 compact	accepted	45.944	8010
862101	Collapsed Taylor	accepted	41.665	7080
862102	Collapsed Taylor	accepted	48.582	8054
862103	Collapsed Taylor	accepted	39.165	6700
862104	Collapsed Taylor	accepted	38.941	6678
862105	Collapsed Taylor	accepted	35.904	6265
862106	Collapsed Taylor	accepted	53.893	8837
862101	Cao ED	accepted	53.866	6788
862102	Cao ED	accepted	66.355	8142
862103	Cao ED	accepted	55.754	6978
862104	Cao ED	accepted	53.004	6559
862105	Cao ED	accepted	49.436	6301
862106	Cao ED	accepted	65.182	8019
862101	Torch coordinate AD	accepted	65.903	6651
862102	Torch coordinate AD	accepted	94.703	9259
862103	Torch coordinate AD	accepted	72.703	7239
862104	Torch coordinate AD	accepted	68.661	6866
862105	Torch coordinate AD	accepted	61.316	6233
862106	Torch coordinate AD	accepted	84.040	8278
862101	STDE exact (2D)	accepted	77.324	6708
862102	STDE exact (2D)	accepted	92.869	7910
862103	STDE exact (2D)	accepted	79.977	6920

Seed	Method	Outcome	Setup + opt.	Closures
862104	STDE exact (2D)	accepted	76.264	6629
862105	STDE exact (2D)	accepted	73.246	6395
862106	STDE exact (2D)	accepted	101.451	8549
862101	PhysicsNeMo	accepted	73.710	6818
862102	PhysicsNeMo	accepted	93.615	8718
862103	PhysicsNeMo	accepted	60.251	6846
862104	PhysicsNeMo	accepted	70.725	6542
862105	PhysicsNeMo	accepted	71.302	6418
862106	PhysicsNeMo	accepted	95.941	8988
862101	cuDNN convolution AD	accepted	516.349	6709
862102	cuDNN convolution AD	accepted	699.179	9049
862103	cuDNN convolution AD	accepted	546.132	7110
862104	cuDNN convolution AD	accepted	498.826	6497
862105	cuDNN convolution AD	accepted	488.845	6367
862106	cuDNN convolution AD	accepted	662.827	8580
862101	cuEquivariance polynomial	accepted	35.119	6654
862102	cuEquivariance polynomial	accepted	45.720	7972
862103	cuEquivariance polynomial	accepted	36.338	6810
862104	cuEquivariance polynomial	accepted	34.258	6482
862105	cuEquivariance polynomial	accepted	33.449	6208
862106	cuEquivariance polynomial	accepted	47.182	8067